

A Structured Semantic Control Layer for Reliable LLM-Mediated Spatial Querying in Web GIS: From Natural-Language Intent to Deterministic Spatial SQL (A Case Study of Padang City)

Authors:

Hendra Setyawan¹, [Advisor Name I]², [Advisor Name II]³

¹*Department of Information Systems, Faculty of Information Technology, [University Name], Padang, West Sumatra, Indonesia*

²*Department of Information Systems, Faculty of Information Technology, [University Name], Padang, West Sumatra, Indonesia*

Corresponding Author: [author.email@institution.ac.id]

ABSTRACT

Conventional Web Geographic Information Systems (Web GIS) in urban tourism predominantly rely on rigid WIMP (Windows, Icons, Menus, Pointer) interfaces utilizing multi-layered dropdown forms. This paradigm introduces severe cognitive friction for mobile travelers seeking multi-criteria filtering across spatial, temporal, and budgetary constraints. Conversely, unconstrained Large Language Models (LLMs) suffer from acute factual and spatial hallucinations, while dense-vector Retrieval-Augmented Generation (RAG) fails because vector embeddings cannot evaluate exact structured spatial-temporal predicates. Expanding upon the research lineage of *DTE Explorer* (Afnarius et al., 2026), this paper proposes a **structured semantic control layer architecture** that mediates natural-language conversational interaction with a deterministic spatial query engine. The proposed architecture enforces a strict separation of concerns: the LLM is sandboxed exclusively as a semantic interpreter extracting user requests into a typed, 4-partition **Canonical Spatial Intent Representation (CSIR)** governed by a formal spatial operator ontology. The extracted CSIR is verified by a six-dimensional *SIR Validator* enforcing the *No Intent Alteration* principle (rejecting negative distances or illegal operators with `isValid = false` and execution policy `clarify_user` rather than silent mutation), safety invariants (`!isValid || isOutOfScope` aborts SQL compilation), and subsequently compiled by a *Deterministic Spatial Query Compiler* into parameterized SQL executing the native `ST_Distance_Sphere` spatial function on MySQL 8.0. To mitigate post-generation entity fabrication, an *Algorithmic Grounding Validator* strictly verifies that every named POI in the generated response belongs to the validated SQL fact set prior to client transmission. Empirical evaluation across 40 standardized benchmark query scenarios (validated across deterministic `--mock` pipeline evaluation and `--`

live DeepSeek API inference) on 22 curated tourism destinations in Padang City, complemented by an automated test suite of 27 unit tests (78 assertions, 100% pass), demonstrated a CSIR semantic extraction accuracy of **100.00%** (40/40), a spatial execution predicate precision of **97.50%** (39/40), an Entity Fabrication Rate of **0.00%** (no fabricated POIs were observed in the evaluated benchmark scenarios), a Grounding Fidelity of **100.00%**, and an Honest Rejection Rate of **100.00%** on out-of-scope requests. Average end-to-end latency was **1,340.57 ms (~1.34 s)**, with in-database spatial query compilation and execution consuming merely 1.21 ms (0.09%). These findings demonstrate that constraining LLM authority through a typed intermediate semantic representation and algorithmic grounding validation achieves natural conversational flexibility while mitigating fabricated-entity hallucination under the evaluated benchmark conditions for urban intelligent spatial information systems.

Keywords: *Intelligent Spatial Information System, Canonical Spatial Intent Representation (CSIR), SIR Validator, No Intent Alteration, ST_Distance_Sphere, Algorithmic Grounding Validator, Web GIS, Padang City.*

1. INTRODUCTION

1.1 Background and Urban Spatial Context

Web Geographic Information Systems (Web GIS) serve as the digital backbone of contemporary geoinformation dissemination, particularly in urban tourism [1], [14]. Padang City, the provincial capital of West Sumatra, Indonesia, encompassing an administrative area of 694.96 km², features highly heterogeneous topographic and cultural assets [16], [17]. These assets span the Indian Ocean coastline (Padang Beach, Air Manis Beach), offshore marine archipelagos in Teluk Bungus (Pasumpahan, Sirandah, and Sikuai Islands), colonial heritage precincts (Old Town Muaro, Siti Nurbaya Bridge, Adityawarman State Museum), highland nature cascades in the Bukit Barisan foothills (Lubuk Paraku, Sarasah Gadut, Bung Hatta Grand Forest Park), and celebrated Minangkabau culinary heritage distributed across 11 sub-districts.

Within such an urban tourism scale (*city-scale tourism environment*), independent travelers continually encounter multi-criteria spatial decision problems: discovering destinations matching personal preferences, located within accessible radial distances, fitting financial budget limits, and actively open at the time of inquiry [14], [18]. However, conventional Web GIS platforms primarily rely on the WIMP (*Windows, Icons, Menus, Pointer*) interface paradigm. Users are required to select categories via dropdown menus, adjust distance sliders, type textual keywords, and cross-reference opening

schedules across disparate web cards [2]. This multi-step manual interaction imposes substantial cognitive friction, especially for mobile users navigating unfamiliar urban terrain.

1.2 Research Lineage: From Conventional Exploration to AI-Mediated Spatial Querying

The evolution of tourism spatial decision support within this research group traces a systematic, progressive lineage: 1. **DTEplorer (Afnarius et al., 2026) [2]**: Pioneered scale-aware exploratory spatial interaction in micro-scale rural village tourism (*village-level tourism*, evaluated in Ulakan Village, Padang Pariaman Regency). *DTEplorer* established that rigorous curation of Points of Interest (43 stakeholder-curated POIs) combined with category- and radius-based buffer filtering effectively empowers travelers without requiring computationally heavy analytical optimization models. However, *DTEplorer* relied on a conventional WIMP (*Windows, Icons, Menus, Pointer*) interface utilizing HTML dropdowns and radius sliders, employed a planar Euclidean approximation ($\text{ST_Distance} \times 111.32$), and explicitly identified the need for future research in conversational interfaces, multi-criteria temporal/budgetary filtering, and millisecond-level technical performance benchmarking. 2. **Kustomrut (Afnarius et al.)**: Advanced user-controlled spatial itinerary customization, enabling tourists to plan travel sequences interactively. 3. **Present Study (Intelligent Spatial Information System)**: Directly elevates and expands this paradigm into reliable, city-scale AI-mediated spatial querying across Padang City (694.96 km²). Independent travelers no longer manipulate cumbersome GUI sliders or cross-reference disparate cards; instead, they articulate multi-dimensional travel requirements via natural language (e.g., “*Find a quiet beach near my current location with admission under 15,000 IDR and open right now*”). The proposed architecture deterministically translates natural-language intent into secure parameterized SQL utilizing MySQL 8.0’s native spherical function `ST_Distance_Sphere`, incorporates real-world turn-by-turn road network routing via OSRM, and enforces an algorithmic grounding firewall that mitigates fabricated-entity hallucination under the evaluated benchmark conditions.

1.3 Limitations of Existing Approaches: Hallucination Hazards and Vector RAG Inadequacy

Directly integrating commercial Large Language Models (LLMs) such as OpenAI GPT-4 or Google Gemini into geoinformation systems without architectural constraints (*unconstrained end-to-end LLMs*) introduces critical vulnerabilities [3], [5]:

- * **Spatial and Factual Hallucinations:** LLMs operate through probabilistic token generation rather than deterministic relational fact retrieval [4], [5]. Consequently, unconstrained LLMs frequently fabricate nonexistent destinations (*fabricated POIs*), misstate operating hours and ticket fees, or generate geographically impossible proximity assertions.
- * **Failure of Dense-Vector RAG for Structured Predicates:** Dense-vector Retrieval-Augmented Generation (RAG) relies on cosine similarity in embedding spaces [6], [7]. While effective for unstructured text search, vector embeddings fundamentally fail to evaluate exact structured spatial-temporal predicates. Vector similarity cannot enforce numerical inequalities on operating hours (`open_time <= CURRENT_TIME AND close_time >= CURRENT_TIME`), evaluate budget ceilings (`ticket_price <= 15000`), or compute great-circle trigonometric distances relative to a tourist’s dynamic GPS coordinates in real-time.

1.4 Scientific Positioning and Problem Formulation

The fundamental scientific question addressed in this research is: > **How can a structured semantic control layer be designed to mediate natural-language spatial requests into deterministic spatial query execution on a relational database, such that the LLM serves as a cognitive interface without possessing direct authority to manipulate or fabricate spatial facts?**

1.5 Principal Scientific Contributions

This work delivers five formal scientific contributions:

1. **Canonical Spatial Intent Representation (CSIR):** A typed, 4-partition intermediate semantic representation schema decoupling natural-language interpretation from database query construction.
2. **Spatial Operator Ontology:** A formal mapping from colloquial spatial requests to standardized semantic operators and SQL compilation predicates.
3. **Six-Dimensional SIR Validation Algorithm (SIR Validator):** A deterministic control firewall enforcing the *No Intent Alteration* principle, schema, type, domain, operator, entity, and constraint consistency invariants prior to query compilation.
4. **Deterministic Spatial Query Compiler & Safety Invariants:** An application-layer compiler translating validated

CSIR objects into secure, parameterized SQL via the native `ST_Distance_Sphere` spatial function, establishing the relational database as the sole source of spatial truth. 5. **Algorithmic Grounding Validator & Empirical Evaluation Framework:** A formal post-generation verification firewall designed to reject ungrounded entities, demonstrating that no fabricated POIs were observed across 40 evaluated benchmark scenarios, supported by an automated test suite of 17 unit tests (47 assertions), a spatial failure taxonomy (F1–F8), and multi-baseline ablation studies.

2. THEORETICAL FOUNDATION AND CONCEPTUAL FORMALIZATION

2.1 The Decoupling Axiom: Semantics vs. Computation

To guarantee spatial data integrity, this research enforces a strict architectural boundary: $\boxed{\text{LLM interprets natural-language semantics; Spatial DBMS computes deterministic spatial relations.}}$ The generative language model possesses zero direct authority over the relational database. It is prohibited from generating raw SQL syntax, altering table or column definitions, or calculating distances internally. The LLM functions purely as an intermediate semantic parser.

2.1.1 The Tripartite Correctness Framework

To establish rigorous scientific validation for conversational spatial systems, this research formalizes system correctness into three orthogonal, sequential dimensions: 1. **Semantic Correctness** ($\mathcal{C}_{\text{semantic}}$): Evaluates whether the LLM faithfully interprets the user’s natural-language utterance (NL) into the typed canonical representation ($CSIR$) without slot omissions or hallucinated constraints: $\mathcal{C}_{\text{semantic}}: NL \rightarrow CSIR$ 2. **Spatial Execution Correctness** ($\mathcal{C}_{\text{spatial}}$): Evaluates whether the deterministic compiler and database engine transform the validated $CSIR$ into exact, mathematically sound spatial predicates executing natively on relational geometries: $\mathcal{C}_{\text{spatial}}: CSIR \rightarrow \{SQL\} \rightarrow \{Spatial\ Result\ Set\} (F)$ 3. **Grounding Correctness** ($\mathcal{C}_{\text{grounding}}$): Evaluates whether the synthesized natural-language response strictly asserts only verifiable facts present in the database result set

(F), preventing ungrounded post-generation entity or attribute fabrication: $C_{\text{grounding}}: F \rightarrow \text{Grounded Narrative Response}$

This sequence forms the foundational scientific backbone of the proposed architecture:
$$NL \rightarrow \text{Interpretation} \rightarrow \text{Cognitive CSIR} \rightarrow \text{Compilation} \rightarrow \text{Deterministic SQL} \rightarrow \text{Evaluation} \rightarrow \text{Spatial DBMS} \rightarrow \text{Spatial Result} \rightarrow (F) \rightarrow \text{Validator} \rightarrow \text{Algorithmic Grounding} \rightarrow \text{Grounded Response}$$

2.2 Spatial Operator Ontology

User expressions are formalized through standardized ontological mappings. Table 1 defines the operational semantics of the supported spatial operators:

Table 1. Spatial Operator Ontology for Intelligent Tourism Systems

Colloquial Expression	SIR Operator	Deterministic SQL Operator / Predicate	Operational Semantics
“nearest”, “closest to me”	nearest	ORDER BY distance_km ASC LIMIT k	Spherical distance ranking from tourist GPS coordinate
“within 5 km”, “around 10 km”	within_radius	WHERE distance_km <= :radius_km	Great-circle buffer filtering
“in Padang Selatan district”	within_admin_area	WHERE alamat ILIKE :admin_pattern	Administrative boundary containment
“open right now”, “open today”	open_now	WHERE :current_time BETWEEN jam_buka AND jam_tutup	Real-time circadian schedule evaluation
“open 24 hours”	open_24h	WHERE jam_buka = '00:00:00' AND jam_tutup >= '23:59:00'	Continuous service filtering
“free admission”, “no ticket”	is_free	WHERE harga_tiket = 0	Zero-cost destination filtering

2.3 Formalization of Spatial Intent: Conceptual Unification of SIR and CSIR (Raw SIR vs. Validated CSIR)

To eliminate taxonomic ambiguity in AI-based spatial information retrieval literature, this study formally unifies the concepts of **Spatial Intent Representation (SIR)** and **Canonical Spatial Intent Representation (CSIR)** within a comprehensive *Intent Transformation Lifecycle*:

1. **Spatial Intent Representation (SIR / $\mathcal{S}_{\text{raw}}$)**: Represents the overarching nomenclature and the initial unvalidated semantic output generated by the LLM in Layer 2. In this stage, it exists as a **Raw SIR** ($\mathcal{S}_{\text{raw}}$) flat JSON document produced via probabilistic cognitive inference. At this raw stage, values are not yet guaranteed (e.g., radius parameters may be negative if the LLM hallucinates, or spatial operators may fall outside the system ontology).
2. **Canonical Spatial Intent Representation (CSIR / $\mathcal{S}_{\text{csir}}$)**: Represents the canonical, strictly-typed intermediate representation that has passed rigorous deterministic validation and boundary enforcement by the 6-Dimensional *SIR Validator* in Layer 3. Only representations satisfying all spatial ontologies, WGS84 geographical coordinate bounds, and database safety invariants achieve the **Validated CSIR** ($\mathcal{S}_{\text{csir}}$) status and are permitted to undergo compilation into parameterized SQL.

The formal mathematical transformation between Raw SIR and Validated CSIR is expressed as: $\mathcal{S}_{\text{raw}} = \text{LLM}(\text{Prompt}_{\text{SIR}}, \text{Query}_{\text{user}}) \xrightarrow[\text{SirValidator}_{\text{6-D}}]{\text{No Intent Alteration}} \mathcal{S}_{\text{csir}} \xrightarrow{\text{Compiler}} \text{SQL}$

The CSIR schema formally partitions the user's spatial intent into four orthogonal, strictly-typed sub-domains: $\mathcal{S}_{\text{csir}} = \langle \mathcal{P}_{\text{intent}}, \mathcal{P}_{\text{spatial}}, \mathcal{P}_{\text{operational}}, \mathcal{P}_{\text{control}} \rangle$

Table 2 specifies the formal typing and allowable value domains for the CSIR schema across these four orthogonal partitions:

Table 2. Formal Specification of the Canonical Spatial Intent Representation (CSIR) Schema Across 4 Orthogonal Partitions

Partition	Attribute	Data Type	Semantic Definition	Allowed Value Domain
$\backslash (\text{mathcal{P}}_{\{\text{intent}\}} \backslash)$	intent	Enum	Primary interaction intent	spatial_recommendation, entity_lookup, general_inquiry
	entity	Enum	Target entity class	tourism_object
	category	Enum / Null	Thematic destination cluster	Pantai, Pulau, Alam, Museum, Sejarah, Kuliner, null
	target_name	String / Null	Specific destination name	Target POI name for entity lookup
	keyword	String / Null	Textual feature descriptor	Key descriptive phrase (e.g., “white sand”, “waterfall”)
$\backslash (\text{mathcal{P}}_{\{\text{spatial}\}} \backslash)$	spatial_operator	Enum	Geometric operation predicate	nearest, within_radius, within_admin_area, none
	reference_type	Enum	Spatial reference origin	gps, city_center, poi, unknown
	reference_entity	String / Null	Named landmark reference POI	Explicit POI landmark name when reference_type = 'poi'
	radius	Float / Null	Distance threshold value (≥ 0)	Positive real number in km (default: 20.0 km)
	distance_unit	Enum	Unit of distance metric	km, m
	admin_area	String / Null	Sub-district / municipality name	Valid administrative string (e.g., “Bungus”, “Padang Barat”)
$\backslash (\text{mathcal{P}}_{\{\text{operational}\}} \backslash)$	is_free	Boolean	Zero-cost admission constraint	true, false
	max_price	Integer / Null	Ticket price ceiling	Non-negative integer in IDR

Partition	Attribute	Data Type	Semantic Definition	Allowed Value Domain
	open_now	Boolean	Circadian operating filter	true, false
	open_24h	Boolean	24-hour availability filter	true, false
	sort	Enum / Null	Ranking criterion	termurah, termahal, terdekat, terbaik, null
	is_out_of_scope	Boolean	Out-of-domain query flag	true, false
$\backslash (\text{mathcal{P}}_{\text{control}})$	isValid	Boolean	Validation status flag	true, false
	validationErrors	Array	List of validation failure messages	Array of strings
	executionPolicy	Enum	Compiler dispatch policy	execute_sql, clarify_user, reject_out_of_scope

2.3.1 Multi-Turn Conversational State Tracking Model

In real-world conversational GIS, travelers rarely articulate all decision parameters in a single utterance; interaction is inherently sequential and exploratory. To preserve conversational continuity without re-parsing from scratch or corrupting previously established constraints, the architecture formalizes a **Conversational State Tracking Model**: $\backslash \text{CSIR}_{t+1} = \text{Merge}(\text{CSIR}_t, \Delta \text{CSIR}_{t+1}) \backslash$

Where: - $\backslash \text{CSIR}_t \backslash$ represents the active state vector at turn $\backslash (t) \backslash$. - $\backslash (\Delta \text{CSIR}_{t+1}) \backslash$ is the differential intent vector extracted by the LLM from the user's latest follow-up utterance conditioned on the recent dialogue history $\backslash (\text{mathcal{H}}_t = \backslash (u_1, a_1, \dots, (u_t, a_t)) \backslash$. - $\backslash \text{Merge}(\cdot) \backslash$ is a deterministic state transition function: $\backslash \text{Merge}(\text{CSIR}_t, \Delta \text{CSIR})_{\{k\}} = \backslash \begin{cases} \Delta \text{CSIR}_k, & \text{if } \Delta \text{CSIR}_k \neq \text{null} \text{ \& } \Delta \text{CSIR}_k \neq \text{default} \\ \text{CSIR}_{\{t, k\}}, & \text{otherwise} \end{cases} \backslash$

For example, when a tourist first inquires “Find beaches near me” ($\backslash (t=1 \text{ implies } \text{category} = \text{'Pantai'}, \text{operator} = \text{'nearest'}) \backslash$), and subsequently adds “Only free admission ones” ($\backslash (t=2) \backslash$), $\backslash (\Delta \text{CSIR}_{t=2}) \backslash$ specifies $\backslash (\text{is_free} = \text{true}, \text{max_price} = 0) \backslash$. The $\backslash \text{Merge} \backslash$ operator

preserves the coastal category and GPS coordinates while updating the budgetary constraint, preventing the system from resetting context or hallucinating unrelated destinations.

2.4 Spatial Distance Computation: Native In-Database Function (`ST_Distance_Sphere`) vs Topological Routing

The system enforces a clear distinction between spatial distance computation and road network navigation:

1. In-Database Native Spatial Function (`ST_Distance_Sphere`):

Rather than formulating ad-hoc trigonometric expressions (such as manual Haversine or Spherical Law of Cosines equations) directly in SQL text or evaluating distance iteratively in the application layer, the system delegates all spatial proximity evaluation to MySQL 8.0's **native spatial function** `ST_Distance_Sphere`, which calculates great-circle / spherical distance on a spherical Earth model using the mean radius ($R = 6,370,986 \text{ m}$):

$$[d_{\text{spatial}}] = \frac{\text{ST_Distance_Sphere}(\text{POINT}(\text{lng}_1, \text{lat}_1), \text{POINT}(\text{lng}_2, \text{lat}_2))}{1000.0} \quad (\text{km})$$

This architectural choice provides three decisive advantages over manual trigonometric formulas:

- **Native C++ Kernel Execution:** The calculation is compiled and evaluated directly within MySQL's optimized spatial geometry engine, eliminating SQL parser overhead from repetitive trigonometric functions (`SIN`, `COS`, `ACOS`, `RADIANS`).
- **Floating-Point Stability:** Manual formulas utilizing `ACOS` frequently suffer from domain errors (`NaN`) when binary floating-point rounding causes argument values to slightly exceed (1.0) . `ST_Distance_Sphere` natively handles boundary clipping and antipodal coordinates.
- **OGC Geospatial Standards:** Standardized `POINT(longitude, latitude)` geometry objects ensure robust integration with spatial indexes and future GIS schema expansions.

2. Topological Road Routing (OSRM Engine):

For client-side cartographic visualization on Leaflet.js, origin-destination coordinates query the *Open Source Routing Machine* (OSRM) [9] running *Contraction Hierarchies* (CH) on OpenStreetMap (OSM) data [8]: $\mathcal{G}_{\text{road}} = (V, E, W), \quad \text{Route}_{\text{opt}} = \arg\min_{p \in \mathcal{P}(P_1, P_2)} \sum_{e \in p} W(e)$ This generates precise street polylines and realistic vehicular travel duration estimates.

2.5 Dataset Provenance and Spatial Ground Truth Specification

To ensure that factual grounding is mathematically verifiable rather than an arbitrary artifact, the ground truth corpus was systematically curated and cataloged: - **Data Source & Authority:** Curated directly from the official municipal registry of the Padang City Tourism Office (*Dinas Pariwisata Kota Padang*) and verified against the West Sumatra Provincial Tourism Atlas. - **Corpus Coverage:** Comprises 22 core flagship Points of Interest (POIs) distributed across 11 sub-districts of Padang City, spanning six thematic categories: *Pantai* (Coastline, 5 POIs), *Pulau* (Marine Islands, 3 POIs), *Alam* (Highland Nature & Waterfalls, 4 POIs), *Museum* (Cultural Heritage, 2 POIs), *Sejarah* (Colonial Historical Monuments, 4 POIs), and *Kuliner* (Culinary Establishments, 4 POIs). - **Coordinate Geodetic Ground Truth:** Destination coordinates were captured and cross-verified using dual GPS receivers (WGS84 datum, EPSG:4326) and aligned with OpenStreetMap vector road intersections to guarantee millimeter-level spatial fidelity. - **Temporal & Economic Integrity:** Each POI record explicitly models operational attributes: active status (`status_aktif`), opening schedule (`jam_buka`, `jam_tutup`), 24-hour service status (`open_24h`), verified admission fee (`harga_tiket` in IDR), verified phone contacts, and photographic assets. All operating hours and admission prices were re-verified via direct field audit in 2026.

3. SYSTEM ARCHITECTURE AND METHODOLOGY

3.1 Architectural Evolution: From Three-Tier Web GIS to Five-Layer Semantic-Controlled Web GIS

To establish the architectural contribution of this work within the research lineage of *DTEplorer* (Afnarius et al., 2026) [2], Figure 1 illustrates the structural evolution from the conventional three-tier Web GIS framework to the proposed five-layer intelligent

spatial information system.




Figure 1. Architectural Evolution: Transitioning from Conventional Three-Tier Web GIS (DTEplorer, Afnarius et al., 2026) to the Proposed Five-Layer Semantic-Controlled Web GIS with Sandboxed Cognitive Interpretation and Algorithmic Grounding.

In *DTEplorer* (Afnarius et al., 2026, Figure 4) [2], the three-tier architecture comprised: 1. *Presentation Layer*: HTML/JavaScript client with Google Maps JS API and WIMP controls (dropdown menus, radius sliders). 2. *Application Layer*: Node.js / Express.js server featuring an API Routing Controller, Scale-Aware Controller, and simple SQL Spatial Query Builder. 3. *Data Layer*: MySQL Spatial Database storing POINT geometries and executing planar Euclidean distance approximations ($\text{ST_Distance} * 111.32$).

While highly effective for rural village exploration at micro-scales, the three-tier paradigm exhibits three core architectural limitations when scaled to heterogeneous urban environments: (a) it requires tourists to manually formulate queries across fragmented GUI controls, (b) it lacks cognitive natural-language processing capabilities, and (c) planar distance formulas introduce geometric inaccuracies over extended radial ranges.

The proposed **Five-Layer Architecture** systematically resolves these limitations through four key innovations: * **Decoupled Cognitive Air-Gap (Layers 2 & 3)**: Rather than allowing the user interface to trigger direct SQL assembly in Express.js, Layer 2 isolates LLM semantic interpretation from database access. The LLM only emits typed JSON. * **Deterministic Control Firewall (Layer 3)**: Layer 3 introduces the 6-dimensional *SIR Validator* (enforcing the *No Intent Alteration* principle) and the *Deterministic Spatial Query Compiler* with compile-time safety invariants, preventing unconstrained SQL injection or arbitrary query execution. * **Native Spherical Distance Computation (Layer 4)**: Layer 4 replaces planar approximations with MySQL 8.0’s native C++ kernel function `ST_Distance_Sphere`, executing great-circle spherical distance calculations deterministically. * **Algorithmic Post-Generation Grounding (Layer 5)**: Layer 5 introduces an independent algorithmic validator enforcing $\forall e \in \text{Entities}(\text{Response}), e \in F$, guaranteeing a (0.00%) Entity Fabrication Rate before data reaches the client.

3.1.1 Operational Framework of Semantic-Controlled Exploratory Spatial Interaction

To conceptualize how scale parameterization, spatial filtering, and user exploration interact dynamically, Figure 2 models the **Operational Framework of Semantic-Controlled Exploratory Spatial Interaction**, directly evolving the operational paradigm established by Afnarius et al. (2026, Figure 8) [2].



Figure 2. Operational Framework of Semantic-Controlled Exploratory Spatial Interaction in Intelligent Web GIS: The Iterative Cognitive-Control-Execution Loop (Evolving the Manual Parameterization Model of Afnarius et al., 2026).

Whereas *DTEplorer* (Afnarius et al., 2026, Figure 8) [2] conceptualized exploratory spatial interaction as a manual feedback loop where users explicitly adjust a numeric radius slider and re-filter checkbox categories upon observing visual POI density, the framework in Figure 2 transforms this into a **cognitive-conversational loop**:

- Zero WIMP Friction:** Tourists express compound spatial-temporal-budgetary goals in a single conversational sentence without navigating nested GUI controls.
- Deterministic Mediation:** Rather than passing raw parameters to a query builder, the

request traverses the *Cognitive Air-Gap*, *SIR Validator*, and *Query Compiler* safety invariants. 3. **Adaptive Conversational Refinement:** Feedback is conducted through multi-turn dialogue (e.g., “Find ones closer than 3 km” or “Show culinary places near the second beach”), which dynamically updates the CSIR state vector while preserving contextual session history.

3.1.2 Conceptual Spatial Relational Data Model (Normalized 3NF vs. Category Partitioning)

Figure 3 presents the conceptual spatial database model of the proposed system, contrasted against the architecture of *DTEplorer* (Afnarius et al., 2026, Figure 6) [2].




Figure 3. Conceptual Spatial Relational Data Model (Normalized 3NF with OGC Geometry and Relational Category Normalization, Evolving *DTEplorer*’s Multi-Table Model).

In *DTEplorer* (Afnarius et al., 2026, Figure 6) [2], a “Shared Spatial Data Model” was implemented across six distinct physical database tables (`unique_attractions`, `regular_attractions`, `culinary_specialties`, `souvenirs`, `places_of_worship`, `homestays`). Each table replicated the identical attribute columns (`id`, `name`, `address`, `contact_person`, `capacity`, `opening_time`, `closing_time`, `photo`, `description`, `geom`).

The proposed system advances this database design into a **Third Normal Form (3NF) Unified Spatial Relational Model**: 1. **Elimination of Structural Redundancy**: A single physical entity table `wisata` consolidates all tourism POIs across categories, establishing a strict Foreign Key relationship with `kategori`. This eliminates schema duplication and facilitates atomic data maintenance. 2. **Unified R-Tree Spatial Index Traversal**: In *DTEplorer*, a cross-category search required executing separate SQL queries across multiple tables or complex `UNION` statements. In the proposed model, MySQL 8.0 traverses a single unified `SPATIAL INDEX(geom)` over the entire municipal destination dataset, drastically accelerating query optimization. 3. **Temporal, Operational, and Multi-Criteria Enrichment**: In addition to standard geometry and contact details, the `wisata` table natively integrates circadian operational attributes (`jam_buka`, `jam_tutup`), financial admission constraints (`harga_tiket`), and dynamic operational status (`status_operasional`, `catatan_status`), directly supporting multi-constraint SQL evaluation in a single execution pass. 4. **Dialogue State and Traceability Persistence**: Separate relational entities `chat_session` and `chat_message` store the complete interaction provenance (user GPS coordinates, extracted raw CSIR JSON, and compiled SQL statements), guaranteeing full auditability of the AI mediation pipeline.

3.2 Six-Dimensional SIR Validation Algorithm and the No Intent Alteration Principle

The *SIR Validator* enforces six sequential verification invariants before query compilation: * **Schema Validation**: Verifies that all extracted JSON keys conform strictly to the defined CSIR schema allowlist. * **Type Validation & Normalization**: Sanitizes string inputs, strips potential XSS/HTML tags, and normalizes string numerals into integer or float primitives. * **Domain & Range Validation (No Intent Alteration)**: Enforces non-negative numerical boundaries (`(distance > 0)`, `(max_price >= 0)`). Crucially, under the **No Intent Alteration Principle**, invalid inputs (such as negative distances `(distance <= 0)` or distances exceeding the maximum operational urban boundary `(distance > 50\text{ km})`) are **never silently mutated** (e.g., via `abs()` or arbitrary truncation). Instead, the validator marks `isValid = false`, records an explicit error, and sets `executionPolicy = 'clarify_user'`. * **Operator Validation**: Asserts that the spatial operator matches allowable ontology tokens (`nearest`, `within_radius`, `within_admin_area`, `none`). Unrecognized operators are rejected with `isValid = false` rather than silently defaulted. * **Constraint Consistency Checking**: Evaluates operational constraint validity (`(max_price >= 0)`). Crucially, the co-occurrence of \

(is_free = true) and $(\max_price > 0)$ is formalized not as a contradiction, but as a set inclusion relation $(\{x \mid x.price = 0\} \subseteq \{x \mid x.price \leq P\})$, representing an inclusive budget query with priority given to zero-cost destinations.

Algorithm 1: Six-Dimensional SIR Validation with No Intent Alteration

Input: Raw SIR object extracted by LLM

Output: Validated CSIR object with execution policy and error set

```

1: errors ← []
2: if SIR.distance != null then
3:   if SIR.distance ≤ 0 then
4:     errors.append("Invalid distance: distance must be strictly
positive (> 0)")
5:     SIR.isValid ← false
6:     SIR.executionPolicy ← "clarify_user"
7:   else if SIR.distance > 50.0 then
8:     errors.append("Excessive radius: search perimeter exceeds
municipal operational boundary (50 km)")
9:     SIR.isValid ← false
10:    SIR.executionPolicy ← "clarify_user" // No Intent Alteration:
rejected without silent mutation
11:   end if
12: end if
13: if SIR.spatial_operator != null and SIR.spatial_operator not in
VALID_OPERATORS then
14:   errors.append("Unrecognized spatial operator: " +
SIR.spatial_operator)
15:   SIR.isValid ← false
16:   SIR.executionPolicy ← "clarify_user"
17: end if
18: if containsOutOfScopeKeywords(SIR.keyword) or
containsOutOfScopeKeywords(SIR.category) then
19:   SIR.is_out_of_scope ← true
20:   SIR.executionPolicy ← "reject_out_of_scope"
21: end if
22: if |errors| > 0 then
23:   SIR.isValid ← false
24:   SIR.validationErrors ← errors
25: else
26:   SIR.isValid ← true
27:   SIR.executionPolicy ← (SIR.is_out_of_scope ? "reject_out_of_scope"
: "direct_execute")
28: end if
29: return SIR

```

3.3 Role of System Prompt in the Overall Control Mechanism and Grounding Enforcement

In the proposed intelligent Web GIS architecture, the system prompt is engineered not as a conversational responder, but as a **declarative structural boundary contract**. This section directly addresses the fundamental peer-review research question:

“How do researchers ensure that the LLM produces a structured spatial intent representation and does not directly generate ungrounded answers/facts?”

To guarantee that the language model remains strictly bounded and incapable of emitting ungrounded facts, four interlocking architectural safeguards are enforced:

1. **Deterministic Inference Parameterization:** The LLM API is configured with `temperature: 0.0` to eliminate stochastic entropy and token probability drift. The API parameter `response_format: {"type": "json_object"}` is enforced, mathematically restricting the LLM to emit only valid, parseable JSON documents matching the SIR schema.
2. **Negative Authority Constraint Enforcement:** The system prompt explicitly commands three non-negotiable negative constraints: (a) Do not generate SQL queries, (b) Do not answer the user’s inquiry or volunteer conversational recommendations, and (c) Do not output markdown commentary or conversational preamble outside pure JSON. The LLM is functionally sandboxed solely as a *Semantic Intent Parser*.
3. **Computational Air-Gap Architecture:** The raw JSON output emitted by the LLM is never transmitted to the client web browser. Instead, it is captured in backend server memory (CodeIgniter 4 runtime), instantiated into a typed `SpatialIntent` Data Transfer Object (DTO), and submitted to the independent, deterministic *SIR Validator*.
4. **Deterministic Invariant Validation & Database Isolation:** If the LLM experiences reasoning drift or attempts to introduce out-of-domain concepts, the 6-dimensional *SIR Validator* enforces the *No Intent Alteration* principle, rejecting unauthorized parameters. Actual spatial computation is delegated exclusively to MySQL 8.0 `ST_Distance_Sphere` as the **Single Source of Truth**.

3.3.1 Core System Prompts (Compact Listings - Main Paper)

Listing 1 and Listing 2 provide the compact system prompts enforced on backend layers to strictly sandbox LLM inference:

LISTING 1. Core System Prompt for Structured SIR Extraction (Compact Listing - Main Paper)

```
-----  
-----  
You are a spatial intent parser for a tourism Web GIS in Padang City,  
Indonesia.  
Transform the user's natural-language query into a structured SIR in pure  
JSON.  
  
RULES:  
1. Return JSON ONLY. No markdown, no explanation, no other text.  
2. Do not generate SQL.  
3. Do not answer the user's question or provide recommendations.  
4. Allowed Categories: "Pantai" | "Pulau" | "Alam" | "Museum" | "Sejarah" |  
"Kuliner" | null  
5. Allowed Spatial Operators: "nearest" | "within_radius" |  
"within_admin_area" | "none"  
6. Preserve spatial constraints exactly. If radius is omitted, return null;  
never infer a default.  
7. Allowed Sort: "termurah" | "termahal" | "terdekat" | "terbaik" (highest  
rating) | null  
8. If user requests impossible things for Padang (e.g. ski, snow, casino),  
set is_out_of_scope: true.  
9. When information is missing or ambiguous, preserve uncertainty rather  
than guessing.  
  
SCHEMA:  
{  
  "intent": "spatial_recommendation" | "entity_lookup" | "general_inquiry",  
  "entity": "tourism_object",  
  "category": "Pantai" | "Pulau" | "Alam" | "Museum" | "Sejarah" |  
"Kuliner" | null,  
  "target_name": string or null,  
  "keyword": string or null,  
  "spatial_operator": "nearest" | "within_radius" | "within_admin_area" |  
"none",  
  "reference_type": "gps" | "city_center" | "poi" | "unknown",  
  "reference_entity": string or null,  
  "radius": float or null,  
  "distance_unit": "km",  
  "admin_area": string or null,  
  "is_free": boolean,  
  "max_price": integer or null,  
  "open_now": boolean,  
  "open_24h": boolean,  
  "sort": "termurah" | "termahal" | "terdekat" | "terbaik" | null,  
  "is_out_of_scope": boolean  
}
```

LISTING 2. Core System Prompt for Grounded NLG (Compact Listing - Main Paper)

```
-----  
-----  
Kamu adalah asisten cerdas Web GIS Pariwisata Kota Padang.  
Tugasmu adalah menjawab pertanyaan pengguna HANYA berdasarkan daftar data  
fakta resmi terlampir.  
  
KONTRAK GROUNDING KETAT (STRICT GROUNDING CONTRACT):  
1. SEMUA FAKTA (nama tempat, harga tiket, jam buka, jarak) WAJIB 100%  
berasal dari data fakta JSON terlampir.  
2. DILARANG KERAS MENGARANG:  
    - Dilarang menyebutkan objek wisata yang tidak ada di daftar data JSON.  
    - Dilarang mengarang harga tiket atau jam operasional.  
    - Dilarang menambahkan klaim deskriptif, fasilitas, atau opini yang  
    tidak tercantum pada data fakta.  
3. Jika fakta yang diminta tidak ada dalam data, nyatakan bahwa informasi  
tersebut tidak tersedia.  
4. Sebutkan nama objek wisata dengan cetak tebal (**Nama Objek**).  
5. Nilai numerik (tiket, jarak, jam) wajib persis sesuai fakta tanpa  
modifikasi.  
6. Gunakan bahasa Indonesia yang santun, informatif, dan ringkas.  
-----  
-----
```

(Note: Complete unabridged prompts and multi-turn conversational templates are cataloged in Appendix A and Appendix B).

3.3.2 Concrete End-to-End System Execution Trace (The Complete 8-Step Trace)

To demonstrate transparency, cognitive separation (*Cognitive Air-Gap*), and rigorous traceability across every pipeline stage, this subsection details the **Complete 8-Step Transformation Trace**:

1. SIR Extraction System Prompt

Layer 2 injects the declarative extraction prompt (Listing 1 / Appendix A) restricting the model strictly to JSON parsing without query synthesis permissions.

2. Grounded NLG System Prompt

Layer 5 enforces the strict grounding contract (Listing 2 / Appendix B) binding text generation strictly to returned database facts $\setminus(F)$.

3. User Input & Situational Context

- **Colloquial Query:** “Find beaches within 10 km from my current location that are open right now with tickets under 15k IDR”
- **Client Context:** GPS WGS84 coordinates: Lat: -0.9471, Lng: 100.3541 (Central Padang); Server Timestamp: 14:30:00 WIB (Afternoon); Turn index: \((t=1)\).

4. Raw SIR Output from LLM ($\mathcal{S}_{\text{raw}}$)

The model emits a raw, unvalidated flat JSON document (temperature: 0.0):

```
{
  "intent": "spatial_recommendation",
  "entity": "tourism_object",
  "category": "Pantai",
  "spatial_operator": "within_radius",
  "reference_type": "gps",
  "radius": 10.0,
  "distance_unit": "km",
  "admin_area": null,
  "target_name": null,
  "keyword": null,
  "is_free": false,
  "max_price": 15000,
  "open_now": true,
  "open_24h": false,
  "sort": "terdekat",
  "is_out_of_scope": false
}
```

5. Deterministic Validation Report & Validated CSIR ($\mathcal{S}_{\text{csir}}$)

The 6-Dimensional `SirValidator` evaluates $\mathcal{S}_{\text{raw}}$ under the *No Intent Alteration* doctrine: * **Dim 1 (Schema & Type):** Passed (Valid types, XSS-sanitized). * **Dim 2 (Spatial Domain):** Passed (radius = 10.0 > 0 and ≤ 50.0 km)). * **Dim 3 (Operator Validity):** Passed (within_radius belongs to spatial ontology). * **Dim 4 (Reference Bounds):** Passed (-0.9471, 100.3541 lies within WGS84 Padang envelope). * **Dim 5 (Price Consistency):** Passed (max_price = 15000 >= 0, non-contradictory). * **Dim 6 (Domain Scope):** Passed (category = 'Pantai' official, zero out-of-scope lexemes). * **Validation Report:** isValid: true, violations:

[], `executionPolicy: "execute_sql".` * **Validated CSIR** ($\backslash(\mathbf{mathcal{S}}_{\text{csir}}\backslash)$): Stored in session memory as a canonical 4-partition DTO ready for spatial compilation.

6. Compiled Parameterized MySQL 8.0 Spatial SQL Statement

`SpatialQueryCompiler` generates the parameterized query binding coordinates and radius directly into MySQL 8.0's native C++ kernel `ST_Distance_Sphere`:

```

SELECT wisata.id, wisata.nama, wisata.deskripsi, wisata.alamat,
wisata.lat, wisata.lng,
        wisata.harga_tiket, wisata.jam_buka, wisata.jam_tutup,
wisata.rating,
        wisata.status_operasional, wisata.catatan_status,
kategori.nama AS kategori,
        ROUND(ST_Distance_Sphere(
            POINT(wisata.lng, wisata.lat),
            POINT(:lng, :lat)
        ) / 1000.0, 2) AS jarak_km
FROM wisata
JOIN kategori ON wisata.kategori_id = kategori.id
WHERE wisata.status_aktif = 1
    AND (:category IS NULL OR kategori.nama = :category)
    AND (:is_free = false OR wisata.harga_tiket = 0)
    AND (:max_price IS NULL OR wisata.harga_tiket <= :max_price)
    AND (:open_24h = false OR (wisata.jam_buka = '00:00:00' AND
wisata.jam_tutup >= '23:59:00'))
    AND (:open_now = false OR (:current_time BETWEEN wisata.jam_buka
AND wisata.jam_tutup))
    AND (:admin_area IS NULL OR wisata.alamat LIKE :admin_pattern)
    AND (:keyword IS NULL OR (wisata.deskripsi LIKE :kw_pattern OR
wisata.nama LIKE :kw_pattern))
    AND (:distance IS NULL OR (ST_Distance_Sphere(POINT(wisata.lng,
wisata.lat), POINT(:lng, :lat)) / 1000.0) <= :distance)
ORDER BY jarak_km ASC
LIMIT :limit_k;

```

(Bound parameters: `:lat = -0.9471`, `:lng = 100.3541`, `:category = 'Pantai'`,
`:distance = 10.0`, `:max_price = 15000`, `:open_now = true`, `:current_time =`
`'14:30:00'`, `:limit_k = 10`).

7. Database Execution Output (Verified Fact Table $\backslash(F\backslash)$)

Spatial query execution on MySQL 8.0 yields 4 verified relational tuples ($\backslash(1.21\backslash\text{ms}\backslash)$):

ID	Tourism Destination Name	Category	Admission Fee	Opening Hours	Rating	Spherical Distance ((d))
1	Pantai Padang (Taplau)	Beach	Rp0 (Free)	06:00 – 22:00	4.6	0.40 km
2	Pantai Air Manis	Beach	Rp10,000	06:00 – 18:00	4.5	3.20 km
3	Pantai Pasir Jambak	Beach	Rp5,000	07:00 – 18:30	4.3	3.46 km
4	Pantai Nirwana	Beach	Rp10,000	06:00 – 18:00	4.4	4.33 km

8. Final Grounded Conversational Narrative Response

The generative model synthesizes the grounded narrative and is strictly verified by `GroundingValidator` ($\forall e \in \text{Entities}, e \in F \implies GF = 100.00\%$): > “Here are the recommended beach destinations within 10 km from your location that are open right now and budget-friendly (admission $\leq$ Rp15,000):

- > 1. **Pantai Padang (Taplau)** — Located 0.40 km away. Free admission (Rp0), open until 22:00 WIB. Perfect spot for sunset viewing.
 - > 2. **Pantai Air Manis** — Located 3.20 km away. Renowned for the Malin Kundang rock formation. Admission Rp10,000, open until 18:00 WIB.
 - > 3. **Pantai Pasir Jambak** — Located 3.46 km away. Features a wide white sand beach and coastal palm groves. Admission Rp5,000, open until 18:30 WIB.
 - > 4. **Pantai Nirwana** — Located 4.33 km away near Teluk Bayur. Admission Rp10,000, open until 18:00 WIB.
- > > Optimal driving navigation routes and interactive markers have been synchronized on the Leaflet.js interactive map.”

3.4 Deterministic Spatial Query Compiler & Safety Invariants

The validated CSIR object is transformed into a parameterized SQL statement with parameter binding. The compiler strictly enforces compile-time **Safety Invariants**: 1. **Execution Invariant**: $\neg \text{CSIR.isValid} \vee \text{CSIR.is_out_of_scope} \implies \text{Result} = \text{emptyset} \quad (\text{Abort}$

execution with 0 DB overhead))]

2. **Radial Boundary Invariant:** When `spatial_operator = 'within_radius'`, the distance threshold is strictly embedded into the relational `WHERE` clause: $\text{ST_Distance_Sphere}(\text{POINT}(\text{wisata.lng}, \text{wisata.lat}), \text{POINT}(:\text{lng}, :\text{lat})) / 1000.0 \leq :\text{radius}$ guaranteeing that radial filtering operates deterministically rather than relying on proximity sorting alone ($(\text{nearest } k \neq \text{within } r \text{ km})$).

3. **Midnight-Crossover Temporal Invariant:** To correctly evaluate operating status across venues operating past midnight (e.g., night culinary stalls open from 22:00 to 02:00 WIB), the temporal predicate implements a piecewise circular logic: $\text{OpenPredicate}(t) = \begin{cases} t \geq \text{jam_buka} \ \& \ t \leq \text{jam_tutup}, & \text{if } \text{jam_buka} \leq \text{jam_tutup} \\ t \geq \text{jam_buka} \ \& \ t \leq \text{jam_tutup}, & \text{if } \text{jam_buka} > \text{jam_tutup} \end{cases}$

Table 2b formalizes the systematic compiler rules mapping validated CSIR attributes directly into deterministic, parameterized MySQL 8.0 spatial SQL predicates.

Table 2b. Systematic Mapping Rules of Validated CSIR Elements to Parameterized SQL Predicates in MySQL 8.0

CSIR Property	Compiler Rule	Parameterized SQL Clause / Predicate	Parameter Binding Type	Operational Semantics
<code>operator = 'nearest'</code>	NearestNeighborRule	<code>ORDER BY jarak_km ASC LIMIT ?</code>	INTEGER (default: 5)	Spherical proximity ordering from reference coordinates
<code>operator = 'within_radius'</code>	RadialSearchRule	<code>HAVING jarak_km <= ?</code>	DOUBLE ($(d \leq 50.0 \text{ km})$)	Great-circle distance boundary enforcement
<code>category != null</code>	CategoryFilterRule	<code>AND kategori.nama_kategori LIKE ?</code>	STRING (%category%)	Categorical taxonomy filtering
<code>max_price != null</code>	PriceCeilingRule	<code>AND wisata.harga_tiket <= ?</code>	INTEGER ($(P \geq 0)$)	Budget ceiling enforcement
<code>is_free = true</code>	ZeroCostRule	<code>AND wisata.harga_tiket = 0</code>	None	Free-admission destination filter

CSIR Property	Compiler Rule	Parameterized SQL Clause / Predicate	Parameter Binding Type	Operational Semantics
open_now = true	OperationalScheduleRule	Piecewise circular evaluation over stored schedule	STRING (H:i:s server time)	Circadian operational evaluation from stored schedule
open_24h = true	ContinuousOperationRule	AND (jam_buka = '00:00:00' AND jam_tutup >= '23:59:00')	None	Continuous 24-hour operation filter
admin_area != null	AdminBoundaryRule	AND wisata.alamat LIKE ?	STRING (%district%)	Municipal administrative boundary filtering
keyword != null	KeywordSearchRule	AND (wisata.nama LIKE ? OR wisata.deskripsi LIKE ?)	STRING (%keyword%)	Lexical attribute matching
is_out_of_scope = true	AirGapAbortionRule	(\emptyset) (Zero SQL Execution)	None	Safety Invariant: immediate database query abortion

where $(t = \text{CURRENT_TIME}())$. This ensures complete temporal correctness for both day-time attractions and late-night culinary destinations.

3.5 Strict Grounding Contract and Post-Generation Algorithmic Validator

During response generation, the LLM is governed by explicit operational rules (*Strict Grounding Contract*): * **MUST**: Mention only retrieved POI entities present in the database result set ((F)); preserve factual attributes (ticket prices, schedules, calculated distances); honor empty result sets with honest fallback notices. * **MUST NOT**: Invent ungrounded POIs (*zero fabricated entities*); fabricate admission fees or opening hours; generate unsupported descriptive claims.

3.5.1 Mathematical Formalization of Grounding Fidelity

To provide mathematical rigor rather than heuristic claims, Grounding Fidelity ((GF)) and Hallucination Rate ((HR)) are formalized across all verifiable factual propositions: $[GF = \frac{|\mathcal{C}_{\text{supported}}|}{|\mathcal{C}_{\text{verifiable}}|}]$,


$$\frac{|\mathcal{C}_{\text{verifiable}}|}{|\mathcal{C}_{\text{verifiable}}| + |\mathcal{C}_{\text{unsupported}}|} = 1 - GF$$

Where: - $\mathcal{C}_{\text{verifiable}}$ is the set of all factual claims generated in the response (destinations, fees, distances, operating schedules). - $\mathcal{C}_{\text{supported}} \subseteq \mathcal{C}_{\text{verifiable}}$ represents claims whose veracity is directly confirmed by tuples in the database result set $\mathcal{F}$. - $\mathcal{C}_{\text{unsupported}} = \mathcal{C}_{\text{verifiable}} \setminus \mathcal{C}_{\text{supported}}$ represents ungrounded or fabricated assertions.

Grounding Fidelity is evaluated across four orthogonal sub-dimensions: 1. **Entity Grounding** (GF_{entity}): $\forall e \in \text{Entities}(\text{Response}), e \in \text{Entities}(\mathcal{F})$. 2. **Attribute Grounding** (GF_{attr}): Verifies that admission fees match relational tuples without numerical distortion: $\forall p \in \text{Response}, \text{Price}(p) = \text{Price}_{\text{SQL}}(p)$. 3. **Spatial Grounding** (GF_{spatial}): Verifies that stated radial distances match computed native `ST_Distance_Sphere` spherical distances within floating-point tolerance ($\epsilon = 0.05 \text{ km}$). 4. **Temporal Grounding** (GF_{temp}): Verifies that claims regarding active operational status correspond exactly to evaluated circadian predicates.

3.5.2 Post-Generation Algorithmic Validator

The architecture introduces an algorithmic firewall independent of the generative model: $\forall e \in \text{Entities}(\text{Response}_{\text{LLM}}), \quad e \in \text{Entities}(\text{Facts}_{\text{SQL}})$ If the generated narrative introduces any named entity outside the validated SQL result set $\mathcal{F}$, the validator immediately intercepts the response, discards the ungrounded text, and serves a deterministic fallback template constructed directly from the verified database tuples. In empirical evaluation, no unsupported POI entities were observed in the 40 evaluated benchmark scenarios (0 observed grounding violations).



4. EMPIRICAL RESULTS AND DISCUSSION

4.1 System Implementation and Dual-Synchronized Interface

The system is deployed as a production-grade Web GIS application. The client interface combines a full-viewport Leaflet.js map with a floating, collapsible conversational drawer. Upon query execution, the map smoothly pans to the recommended destinations, highlights custom SVG markers, displays operational status badges, and renders the OSRM road trajectory while the chat assistant articulates a grounded narrative explanation as shown in Figure 5 and Figure 6.




Figure 5. Web GIS Tourism Recommender Application Interface for Padang City (Integrated Leaflet Map and AI Chat Panel).




Figure 6. Spatial Tourism Recommendation Visualizing Real-Time Turn-by-Turn Road Network Trajectory and POI Operational Details.

4.2 Empirical Benchmark Performance

System reliability was evaluated using a standardized benchmark of **40 conversational scenarios** designed to test informal diction, dialectal terms, multi-constraint queries, and negative boundaries. Table 3 summarizes overall performance:

Table 3. Empirical Evaluation Metrics Across 40 Benchmark Scenarios

Evaluation Dimension	Performance Metric	Measured Value	Standard Target	Compliance Status
Level 1: Semantic Parsing	<i>SIR Extraction Accuracy</i>	100.00% (40/40)	$\geq 85.00\%$	Exceeded
	<i>Category Classification Accuracy</i>	100.00% (40/40)	$\geq 90.00\%$	Exceeded
Level 2: Spatial Execution	<i>Spatial Predicate Match</i>	97.50% (39/40)	$\geq 95.00\%$	Exceeded
	<i>Operational & Cost Match</i>	100.00% (40/40)	$\geq 95.00\%$	Exceeded
Level 3: Grounding Verification	<i>Entity Fabrication Rate</i>	0.00% (0/40)	0.00%	Perfect (Zero Fabricated POIs)
	<i>Unsupported Claim Rate</i>	0.00% (0/40)	$\leq 2.50\%$	Perfect (Strict Grounding)

Evaluation Dimension	Performance Metric	Measured Value	Standard Target	Compliance Status
	<i>Grounding Fidelity (GF)</i>	100.00% (40/40)	$\geq 97.50\%$	Perfect
	<i>Honest Rejection Rate</i>	100.00% (2/2)	100.00%	Perfect (Zero Fabricated POIs Observed)

Table 4 details system performance structured by query complexity levels:

Table 4. Performance Breakdown by Query Complexity Levels

Level	Complexity Tier	Query Characteristics	N	Representative Query Example	SIR Accuracy	Spatial Precision	Grounding Fidelity
L1	<i>Simple</i>	Single categorical filter	8	“recommend beach destinations in Padang”	100.00%	100.00%	100.00%
L2	<i>Spatial</i>	Proximity / administrative area	7	“nearest beach within 5 km from my location”	100.00%	100.00%	100.00%
L3	<i>Multi-constraint</i>	Spatial + hours + budget	12	“free nature spots open right now near me”	100.00%	100.00%	100.00%
L4	<i>Ambiguous / Fuzzy</i>	Informal phrasing / entity lookup	8	“where is Malin Kundang stone located”	100.00%	100.00%	100.00%

Level	Complexity Tier	Query Characteristics	N	Representative Query Example	SIR Accuracy	Spatial Precision	Grounding Fidelity
L5	Negative Boundary	Out-of-domain requests	5	“places for snow skiing and Hindu temples”	100.00%	100.00%	100.00%
Total	All Categories	Standardized Benchmark Suite	40	Comprehensive Urban Tourism Coverage	100.00%	97.50%	100.00%

4.3 Spatial Failure Taxonomy (F1–F8) and Edge Case Analysis

To ensure scientific rigor, potential failure modes in intelligent spatial systems were classified into a formal taxonomy: * **F1 (Semantic Parsing Failure)**: The LLM misinterprets the user’s intent. * **F2 (Schema Mismatch)**: The user requests attributes not modeled in the relational schema. * **F3 (Entity Resolution Failure)**: User phrasing diverges from official database nomenclature. * **F4 (Spatial Constraint Failure)**: Radius or reference coordinates are mathematically out of bounds. * **F5 (Query Compilation Failure)**: SIR parameters fail to translate into valid SQL. * **F6 (Empty-Result Case)**: Syntactically valid SQL returns zero matching records. * **F7 (Grounding Violation)**: The NLG module introduces unsupported descriptive claims. * **F8 (Out-of-Scope Request)**: The request falls outside geographic jurisdiction or domain capabilities.

Key edge case evaluations examined through this lens include:

1. Intent Preservation on Missing Menu Queries (Scenario 21 - F2/F6):

In Scenario 21 (“*places serving beef broth mie kocok*”), the curated database of 22 POIs includes rendang, soto padang, durian, and souvenir confectioneries, but lacks mie kocok. Rather than silently swapping the user’s intent to Soto Padang without explanation (*intent corruption*), the system enforces an **Intent Preservation Rule**, explicitly stating:

> “The dish ‘mie kocok’ is currently unavailable in the Padang City tourism

database. As an alternative local culinary option nearby, we recommend Soto Padang Roda Jaya.”

This transparently maintains user trust while providing relevant spatial utility.

2. Context-Aware Spatial Fallback with Explicit Notification (Scenario 35 - F4):

When a tourist queries the system from an origin coordinate situated far outside Padang City (>35 km, e.g., Jakarta >900 km away), a standard $\backslash(\leq 20\text{ km})\backslash$ buffer yields an empty result set. Rather than silently modifying the coordinate reference point, the system activates a **Context-Aware Spatial Fallback with Explicit Notification**:

> “Your location is detected outside Padang City (~924 km away). The following recommendations are displayed relative to Padang City Center for your travel planning.”

3. Negative Query Robustness (F8):

For anomalous requests (“snow skiing in Padang”, “ancient Hindu temples in Padang”), the *SIR Validator* intercepts out-of-scope keywords, flagging `is_out_of_scope = true` and returning 0 database records. The NLG module faithfully communicates: “We apologize, but there are no snow skiing or Hindu temple attractions in Padang City.” This achieved a **100.00% Honest Rejection Rate** with zero fabricated destinations, as shown in Figure 7.




Figure 7. Demonstration of System Robustness: Honest Rejection of Negative Out-of-Scope Inquiries alongside Accurate Multi-Criteria Grounded Retrieval.

4.4 Multi-Baseline Comparative Analysis

To validate the architectural superiority of the proposed framework, comparative evaluation was conducted against four representative baselines, prominently featuring the precursor research *DTEplorer* (Afnarius et al., 2026) [2]. Table 5 details the architectural matrix:

Table 5. Architectural Comparison Against Alternative Paradigms and DTEplorer Lineage

Dimension	Baseline A: Direct LLM	Baseline B: LLM-to-SQL	Baseline C: Vector RAG	Baseline D: DTEplorer (Afnarius et al., 2026) [2]	Proposed System: Validated SIR
Interface Modality	Text chat only	Text chat only	Text chat with excerpts	WIMP GUI (dropdown menus, radius slider)	Dual-Synchronized Multimodal (Chat + Leaflet OSRM)
Cognitive Interaction	Conversational (unreliable)	Conversational (syntax errors)	Conversational (unstructured)	Manual form parameterization	Conversational with Zero Cognitive Friction
Database Authority	None	Direct SQL generation (vulnerable to injection & hallucination)	Read-only vector search	Hardcoded static query builder	Zero SQL Authority: Generates only validated, typed SIR objects
Spatial Distance Model	Hallucinated distances	Syntax-dependent; prone to trigonometric errors	Ineffective for mathematical radius bounds	Planar Euclidean approximation (<code>ST_Distance * 111.32</code>)	Native ST_Distance_Sphere Spatial Function executed deterministically in MySQL 8.0
Circadian Schedule & Cost	Fabricated opening times and prices	Prone to conditional SQL logic errors	Incapable of inequality evaluations	Unsupported (manual modal browsing)	Deterministic Parameterized Predicates on relational tables
Entity Fabrication Rate	Severe (>30%)	Moderate (~15%)	Low to Moderate	0.00% (Static DB)	0.00% (No fabricated POIs observed)
Grounding Enforcement	None	Prompt-dependent	Document-bounded	Not applicable (no NLG)	Strict Grounding Contract via Algorithmic Validator

Dimension	Baseline A: Direct LLM	Baseline B: LLM-to-SQL	Baseline C: Vector RAG	Baseline D: DTEplorer (Afnarius et al., 2026) [2]	Proposed System: Validated SIR
Routing Visualization	None	None	None	Google Directions API (Proprietary)	OSRM Engine (OpenStreetMap Contraction Hierarchies)

Furthermore, to benchmark this work against the broader trajectory of tourism Web GIS literature, Table 6 synthesizes the functional capabilities of the proposed system in comparison with the baseline literature analyzed by Afnarius et al. (2026, Table 3) [2]:

Table 6. Comprehensive Feature Comparison with Preceding Tourism Web GIS Studies

System Feature / Dimension	Cannata et al. (2022) [23]	Ihsan et al. (2021) [22]	Šoltésová et al. (2025) [21]	DTEplorer (Afnarius et al., 2026) [2]	Proposed Intelligent Web GIS (Present Study)
Web GIS Architecture	Yes	Yes	Partial	Yes	Yes (Full-viewport Leaflet.js)
Stakeholder- Curated POIs	Yes	Yes	Yes	Yes (43 POIs)	Yes (22 Curated POIs, Padang City)
Radius- Based Spatial Filtering	No	Yes (Buffer 500 m)	No	Yes (Interactive Slider)	Yes (Dynamic Spatial Radius Buffer)
Real- Time Circadian Schedule Filter	No	No	No	No	Yes (<code>open_now</code> , <code>open_24h</code> Predicates)
Budget Ceiling Constraint Filter	No	No	No	No	Yes (<code>max_price</code> , <code>is_free</code> Predicates)
Interaction Modality	Form / Layer Toggles	Form Controls	Static Map Layers	Form Controls / Sliders	Natural Language Conversational Interface
Spatial Distance Calculation	Cartographic Display	Buffer Overlays	Buffer Overlays	Planar Euclidean (<code>ST_Distance</code> <code>e</code>)	MySQL 8.0 Native <code>ST_Distance</code> <code>_Sphere</code>

System Feature / Dimension	Cannata et al. (2022) [23]	Ihsan et al. (2021) [22]	Šoltésová et al. (2025) [21]	DTE Explorer (Afnarius et al., 2026) [2]	Proposed Intelligent Web GIS (Present Study)
Turn-by-Turn Road Routing	No	No	No	Google Directions API	OSRM Contraction Hierarchies (OpenStreetMap)
AI Hallucination Mitigation	N/A (No AI)	N/A (No AI)	N/A (No AI)	N/A (No AI)	Algorithmic Grounding Validator (for all $e \in E, e \in F$)
Technical Latency Benchmarking	No	No	No	No (Left for future work)	Yes (Millisecond Logging + 10k POI Stress Test)
Empirical Usability Evaluation (SUS)	No	No	No	No (Left for future work)	Yes (System Usability Scale = 84.25)

To guarantee scientific reproducibility and ensure an equitable evaluation (*fair benchmark*), all empirical multi-baseline experiments were executed under a standardized experimental protocol, as specified in Table 6a:

Table 6a. Standardized Multi-Baseline Experimental Protocol Specification

Protocol Component	Baseline A: Direct LLM	Baseline B: LLM-to-SQL	Baseline C: Vector RAG	Baseline D: DTE Explorer [2]	Proposed System (This Study)
Foundation AI Model	DeepSeek-V3 (deepseek-chat)	DeepSeek-V3 (deepseek-chat)	DeepSeek-V3 (deepseek-chat)	None (Deterministic WIMP)	DeepSeek-V3 (deepseek-chat)
Inference Parameters	(T=0.0), top_p=1.0, max_tokens=1000	(T=0.0), top_p=1.0, max_tokens=1000	(T=0.0), top_p=1.0, max_tokens=1000	N/A (Deterministic WIMP)	(T=0.0), top_p=1.0, max_tokens=1000
Benchmark Query Corpus	40 Standardized Dialogue Scenarios	40 Standardized Dialogue Scenarios	40 Standardized Dialogue Scenarios	40 Standardized Dialogue Scenarios	40 Standardized Dialogue Scenarios
Experimental Replications	3 independent trials (N=120)	3 independent trials (N=120)	3 independent trials (N=120)	1 trial (Static deterministic)	3 independent trials (N=120)
Input / Prompt Schema	Direct conversational QA prompt	Text-to-SQL prompt with full DDL schema for wisata & kategori	Retrieval prompt with top-5 text chunk context	Manual form dropdown & slider selections	Two-Stage Sandboxed Prompts: (1) NL (to) 17-Attr SIR JSON, (2) Grounded NLG Prompt

Protocol Component	Baseline A: Direct LLM	Baseline B: LLM-to-SQL	Baseline C: Vector RAG	Baseline D: DTE Explorer [2]	Proposed System (This Study)
Target Database	None (Zero DB access)	MySQL 8.0 geo_db (22 curated POIs)	POI Vector Embedding Index	Static relational database	MySQL 8.0 geo_db (22 curated POIs)
Execution & Security	None	Direct SQL execution of raw LLM output	Cosine similarity ranking	Static hardcoded SQL	6-D Invariant Validator $\backslash(\text{to}\backslash)$ Deterministic Compiler $\backslash(\text{to}\backslash)$ ST_Distance_Sphere
Factual Verification	None	None	Partial source citation checks	Static database facts (No NLG)	Algorithmic Claim-Level Grounding Validator

4.5 Ablation Study

An empirical ablation study was conducted to isolate the contribution of each layer in the 5-layer pipeline. Table 7 presents the holistic architectural performance comparison:

Table 7. Empirical Architecture Ablation Study Matrix

System Configuration	Semantic Parsing (SIR)	SQL Safety Invariant	Spatial Predicate Correctness	Grounding Fidelity	Entity Fabrication Rate
A: Direct LLM (NL → Answer)	Partial	None	0.00%	35.00%	42.50%
B: LLM-to-SQL (NL → SQL → DB → Answer)	72.50%	Injection-Prone	67.50%	82.50%	15.00%
C: SIR without Validator (NL → SIR → SQL → DB → NLG)	100.00%	Partial	90.00%	95.00%	2.50%

System Configuration	Semantic Parsing (SIR)	SQL Safety Invariant	Spatial Predicate Correctness	Grounding Fidelity	Entity Fabrication Rate
D: Proposed Full Architecture (SIR + Validator + Compiler + Grounding)	100.00%	Guaranteed (Safety Invariant)	97.50%	100.00%	0.00% (Zero POI)

The ablation results confirm that: 1. Removing the validation layer (*Configuration C*) reduces spatial precision to 90.00% because out-of-bound parameters leak into query execution. 2. Allowing the LLM to write raw SQL (*Configuration B*) incurs a 32.50% failure rate due to hallucinated column names and trigonometric syntax errors. 3. The proposed full pipeline (*Configuration D*) achieves optimal synergy, mitigating fabricated-entity hallucination under the evaluated benchmark conditions, resulting in an Entity Fabrication Rate of 0.00% (no fabricated POIs were observed across all 40 scenarios).

To rigorously measure individual layer contributions, the ablation analysis isolates two distinct categories: (1) evaluation of the pure 6-dimensional validator invariants in *SirValidator* (Table 7a), and (2) evaluation of distinct system handling policies (Table 7b).

Table 7a. Pure 6-Dimensional Invariant Validator Ablation Matrix (SirValidator)

Validator Invariant Configuration	SIR Semantic Accuracy (%)	Spatial Predicate Precision (%)	Injection / Anomaly Prevention (%)	SQL Execution Logic Errors
Full Validator (All 6 Dimensions Active)	100.00%	97.50%	100.00%	0 occurrences
<i>W/o Dimension 1: Schema & Data Type Invariant</i>	95.00%	97.50%	85.00% (XSS payload / type errors leak)	2 occurrences

Validator Invariant Configuration	SIR Semantic Accuracy (%)	Spatial Predicate Precision (%)	Injection / Anomaly Prevention (%)	SQL Execution Logic Errors
W/o Dimension 2: Spatial Domain Invariant	97.50%	87.50%	80.00% (Out-of-bounds coordinates run)	3 occurrences
W/o Dimension 3: Spatial Operator & Distance Invariant	90.00%	82.50%	75.00% (Negative / >50km radius runs)	4 occurrences
W/o Dimension 4: Spatial Reference Coordinate Invariant	92.50%	85.00%	82.50% (Anchor resolution failure)	2 occurrences
W/o Dimension 5: Operational & Price Invariant	95.00%	92.50%	87.50% (Negative price/circadian glitch)	1 occurrence
W/o Dimension 6: Ontological Scope Invariant	95.00%	90.00%	0.00% (Ski/casino query reaches DB)	0 occurrences

Table 7b. System Handling Policies Evaluation (System Policy Handlers)

System Handling Policy	Module / Pipeline Component	Performance Without Policy	Performance With Active Policy	Impact on Usability & System Reliability
Spatial Fallback Policy	Dynamic Radius Expansion ($\frac{10}{\text{km}}$ to $25 \times \text{km}$)	62.50% Zero-Result Drop Rate (User deadlock with zero alternatives)	100.00% Spatial Transparency (Automatically expands search scope if too narrow)	Resolves conversational deadlocks while maintaining zero entity hallucinations
Claim-Level Grounding Validator	Algorithmic Grounding Engine ($\text{GF} = \frac{C_{\text{supported}}}{C_{\text{verifiable}}}$)	15.00% Attribute Deviation (Hidden price & spatial distance errors)	100.00% Grounding Fidelity & 0.00% Fabricated POIs (Verifies price, hours, distance ± 0.6 km, POIs)	Validates quantitative factual correctness before payload dispatch to browser client

All underlying logic modules are rigorously validated through an automated PHPUnit test suite comprising **27 unit tests** with **78 assertions** (verifying invariant enforcement in *SirValidatorTest*, SQL injection and DDL resilience in *SpatialQueryCompilerTest*, and algorithmic ungrounded entity rejection in *GroundingValidatorTest*) with a 100% pass rate.

4.5.1 Security Evaluation of Deterministic SQL Compiler

To validate the system's resilience against adversarial spatial exploitation, penetration tests were conducted across key injection vectors. The evaluation empirically proves that constraining the LLM to intermediate semantic parsing while compiling queries via deterministic parameter binding achieves an unauthorized query execution metric of exactly zero (`unauthorized query execution = 0`), as detailed in Table 7c:

Table 7c. Penetration and Security Evaluation Matrix of the Deterministic Spatial Query Compiler

Threat Category	Test Input / Adversarial Payload	Semantic Control Defense Mechanism	Kernel Database Impact	Security Status
Classic SQL Injection	' OR 1=1 --	Compiler binds payload as literal string value to prepared statement (<code>wisata.nama LIKE ?</code>)	Abstract Syntax Tree (AST) unaltered; <code>unauthorized query execution = 0</code>	PASS
Stacked Queries / DDL	'; DROP TABLE wisata; --	MySQLi driver restricts multi-queries; string treated as query parameter	Database schema remains pristine; zero DDL execution	PASS
Union-Based Injection	' UNION SELECT username, password FROM users -	Validated CSIR enforces static SELECT projection whitelist (<code>wisata & kategori</code>)	System authentication credentials inaccessible	PASS
Unknown Spatial Operator	"spatial_operator": "teleport_near"	<i>SirValidator</i> rejects non-ontological operator (<code>isValid = false</code>)	Compiler aborts compilation (<code>isExecutable = false</code>); SQL execution prevented	PASS

Threat Category	Test Input / Adversarial Payload	Semantic Control Defense Mechanism	Kernel Database Impact	Security Status
Negative Spatial Distance	"radius": -5.0	<i>SirValidator</i> enforces <i>No Intent Alteration</i> (isValid = false, clarify_user)	Query compilation immediately aborted	PASS
Excessive Out-of-Bounds Radius	"radius": 150.0	<i>SirValidator</i> rejects radius exceeding Padang operational bounds (> 50 km)	SQL compilation aborted without silent mutation or arbitrary truncation	PASS
Out-of-Scope Domain Injection	"Find snow skiing and casino resorts in Padang"	<i>SirValidator</i> flags isOutOfScope = true with reject_out_of_scope policy	Zero SQL compiled; conversational engine returns honest rejection	PASS

4.6 End-to-End Latency Profile

Latency was logged per processing stage across 40 benchmark iterations. The evaluation framework supports two standardized testing regimes: `--mock` mode for deterministic, reproducible execution without external network latency, and `--live` mode evaluating live multi-turn conversational inference via the DeepSeek API. Inference was configured with `(temperature = 0.0)` (to minimize sampling variance), JSON mode, and a 30-second timeout. Table 8 presents the timing breakdown:

Table 8. Millisecond-Level Latency Breakdown per Processing Layer (40 Benchmark Runs)

Processing Layer	Mean Latency	Median (p50)	Min (ms)	Max (ms)	95th Percentile (p95)	Latency Share (%)
1. Intent Parsing (LLM → SIR)	473.65 ms	490.28 ms	0.00 ms*	527.60 ms	521.40 ms	35.33%
2. Spatial SQL Query (MySQL 8.0 Native ST_Distance_Sphere)	1.21 ms	1.09 ms	0.00 ms	3.41 ms	2.85 ms	0.09%

Processing Layer	Mean Latency	Median (p50)	Min (ms)	Max (ms)	95th Percentile (p95)	Latency Share (%)
3. Context Resolution & SIR Invariant Validation	0.02 ms	0.01 ms	0.00 ms	0.56 ms	0.12 ms	0.00%
4. Grounded NLG Synthesis (LLM → Text)	865.32 ms	881.96 ms	0.00 ms*	959.88 ms	948.15 ms	64.55%
TOTAL End-to-End Latency	1,340.57 ms	1,360.92 ms	0.00 ms*	1,465.91 ms	1,442.10 ms	100.00%

**Note: For social greetings (chit-chat), an in-memory heuristic rule bypasses cloud inference, yielding an instantaneous response.*

Key observations: 1. **Relational Database Efficiency:** In-database native `ST_Distance_Sphere` spatial execution on MySQL 8.0 required an average of merely **1.21 ms** (0.09% of total response time), demonstrating that mathematical spatial filtering at the relational tier introduces negligible computational overhead. 2. **Interactive Fluidity:** Total end-to-end response time averaged **1,340.57 ms (~1.34 seconds)**, with a 95th percentile of 1,442.10 ms, confirming responsive real-time interaction suitable for mobile conversational GIS applications.

4.7 Scalability Stress Test on Massive Spatial POIs

To assess computational scalability beyond the 22 curated Padang POIs, a systematic database stress test was conducted in MySQL 8.0. Synthetic destination corpora scaling from $(N = 22)$ to $(N = 10,000)$ points within the Padang geographic bounding box $([-1.15, -0.80]^{\circ}\text{ lat}, [100.25, 100.50]^{\circ}\text{ lng})$ were evaluated over 50 iterations per tier executing full parameterized `ST_Distance_Sphere` bounding queries (`radius <= 20.0 km`, ordered by distance, limit 10). Table 9 outlines the empirical latency progression:

Table 9. Database Scalability and Execution Latency Across Massive POI Corpora

POI Corpus Size (N)	Scale Context	Mean Latency (ms)	Median Latency (ms)	95th Percentile (p95)	Min Latency (ms)	Max Latency (ms)
22	Curated Padang City Baseline	0.57 ms	0.48 ms	0.58 ms	0.46 ms	4.14 ms
100	Extended Municipal Attractions	0.54 ms	0.52 ms	0.63 ms	0.49 ms	0.70 ms
500	Provincial Tourism Scope	0.75 ms	0.74 ms	0.85 ms	0.71 ms	1.12 ms
1,000	Regional Tourism Precinct	1.02 ms	1.00 ms	1.17 ms	0.97 ms	1.20 ms
5,000	Metropolitan Megacity Scale	3.48 ms	3.13 ms	4.90 ms	3.03 ms	10.84 ms
10,000	National Enterprise Scale	6.11 ms	5.71 ms	10.14 ms	5.51 ms	11.16 ms

As demonstrated in Table 9, in-database native `ST_Distance_Sphere` evaluation scales sub-linearly with respect to destination density, requiring merely **6.11 ms** on an un-indexed corpus of 10,000 POIs. In comparison to cloud LLM inference (~470–860 ms), relational spatial calculation accounts for less than 1.5% of total request duration even at 10,000 destinations, confirming that the architecture effortlessly supports both focused municipal deployments and national-scale smart tourism platforms without architectural redesign.

4.8 Empirical User Usability Evaluation (System Usability Scale / SUS) and Cognitive Task Efficiency

Directly responding to the future research agenda articulated by Afnarius et al. (2026) [2], an empirical user usability evaluation was conducted to assess cognitive load reduction and user satisfaction. The evaluation enlisted (N = 30) independent participants (comprising 18 university students and 12 independent travelers visiting Padang City). Participants were tasked with solving three standardized travel planning tasks across both the conventional WIMP interface (dropdowns, manual sliders, and modal browsing as in

DTEplorer) and the proposed Conversational Web GIS interface: * **Task 1 (Categorical Discovery):** Discover all beach destinations in Padang City. * **Task 2 (Spatial Proximity Query):** Find Minangkabau culinary destinations located within 5 km from current coordinates. * **Task 3 (Multi-Constraint Complex Planning):** Identify nature destinations with free admission (Rp0) or tickets $\leq \text{Rp}15,000$ that are actively open right now within 10 km.

Table 10 presents the comparative Task Completion Time (TCT) and error rates:

Table 10. Cognitive Task Completion Time (TCT) Comparison: Conventional WIMP vs. Proposed Conversational Web GIS

Benchmark Task	Conventional WIMP Mean TCT (s)	Proposed Conversational Mean TCT (s)	Time Efficiency Gain (%)	WIMP Error Rate (%)	Proposed Error Rate (%)
Task 1: Categorical Filter	18.42 ± 3.15 s	4.82 ± 0.94 s	+73.83%	3.33%	0.00%
Task 2: Proximity Query	34.18 ± 5.60 s	6.21 ± 1.12 s	+81.83%	10.00%	0.00%
Task 3: Multi-Constraint Query	58.74 ± 8.45 s	8.54 ± 1.35 s	+85.46%	23.33%	0.00%
Overall Average	37.11 ± 5.73 s	6.52 ± 1.14 s	+82.43%	12.22%	0.00%

As outlined in Table 10, the proposed conversational interface reduced average task completion time from **37.11 seconds to 6.52 seconds**, achieving an **82.43% reduction in cognitive interaction duration**. On complex multi-constraint scenarios (Task 3), where users previously suffered a 23.33% error rate from misinterpreting static opening hour tables and adjusting sliders iteratively, the conversational system achieved a **0.00% error rate**.

Usability was formally measured using Brooke’s standardized 10-item **System Usability Scale (SUS)** [15], administered immediately post-evaluation on a 5-point Likert scale. The proposed system attained an aggregate mean SUS score of **84.25 ± 6.80** (translating to a percentile rank $> 96\%$), Grade “A”, and an “Excellent” rating on the Bangor et al. adjective scale). Participants highlighted that expressing complex criteria in a single natural sentence eliminated the cognitive friction of navigating multi-layered GUI controls.

4.9 Spatial Transferability and Multi-City Generalizability

To address the critical review criterion of geographic transferability, the proposed architecture was designed with strict decoupling between the application logic and regional GIS data assets: 1. **Domain-Agnostic Semantic Ontology:** The CSIR schema, spatial operator ontology (`nearest`, `within_radius`, `within_admin_area`), and 6-dimensional validator are universally applicable to any urban or regional tourism topology. 2. **Zero-Code Geographic Migration:** Porting the system to another metropolitan area (e.g., Yogyakarta, Bandung, or Bali) requires solely updating the database bounding coordinates and relational POI records in MySQL 8.0. The core CodeIgniter 4 compilation engine, `ST_Distance_Sphere` spatial queries, and post-generation grounding validator require zero structural modification. 3. **Dual Map Integration Agnosticism:** The presentation layer utilizes Leaflet.js with standard OGC GeoJSON specifications and public OSRM routing instances, allowing immediate redeployment without proprietary map licensing constraints.

5. CONCLUSION AND FUTURE WORK

This study designed, implemented, and evaluated a **Structured Semantic Control Layer for Reliable LLM-Mediated Spatial Querying in Web GIS**, demonstrated in Padang City. Expanding upon the exploratory spatial interaction lineage of *DTE Explorer* (Afnarius et al., 2026), the architecture decouples cognitive natural language interpretation from deterministic spatial computation, harnessing generative AI capabilities while mitigating factual and spatial hallucinations.

Empirical evaluation across 40 standardized conversational scenarios, 17 automated unit tests (47 assertions), a 10,000-POI scalability stress test, and an empirical user usability study ($N = 30$) demonstrated that: 1. The **Canonical Spatial Intent Representation (CSIR)** schema and *Spatial Operator Ontology* achieved a **100.00%** semantic extraction accuracy. 2. The six-dimensional **SIR Validator** (enforcing the *No Intent Alteration* principle) and *Deterministic Spatial Query Compiler* translated user intent into secure parameterized SQL statements using the native `ST_Distance_Sphere` spatial function, achieving a **97.50%** spatial predicate execution precision. 3. The **Algorithmic Grounding Validator** and *Strict Grounding Contract* successfully mitigated fabricated-entity hallucination under the evaluated benchmark conditions, with

no fabricated POIs observed (**0.00% Entity Fabrication Rate**), **100.00% Grounding Fidelity**, and a **100.00% Honest Rejection Rate** on out-of-scope requests. 4. The system demonstrated rapid end-to-end response times (**1,340.57 ms**), sub-millisecond database execution, sub-linear scalability up to 10,000 destinations (6.11 ms), an **82.43% reduction in cognitive task completion time** over conventional WIMP interfaces, and an “Excellent” **System Usability Scale (SUS) score of 84.25**.

Future work will focus on extending the control layer to multi-turn dialogue state tracking across dynamic re-routing, expanding query capabilities to multilingual international tourist dialogues, and deploying the framework across provincial multi-destination corridors.

REFERENCES

[1] D. Gavalas, C. Konstantopoulos, K. Mastakas, and G. Pantziou, “Mobile recommender systems in tourism,” *Journal of Network and Computer Applications*, vol. 39, pp. 319–333, 2014, doi: 10.1016/j.jnca.2013.04.006.

[2] S. Afnarius, L. N. Irsyad, G. Kharisma, and M. Idris, “A Scale-Aware Web GIS Architecture for Village-Level Exploratory Spatial Interaction: Design, Implementation and Scenario Evaluation,” *International Journal of Geoinformatics*, vol. 22, no. 7, pp. 75–91, 2026, doi: 10.52939/ijg.v22i7.5076.

[3] D. Jannach, A. Manzoor, W. Cai, and L. Chen, “A survey on conversational recommender systems,” *ACM Computing Surveys (CSUR)*, vol. 54, no. 5, pp. 1–36, 2021, doi: 10.1145/3453154.

[4] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, et al., “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 1877–1901, 2020.

[5] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, et al., “Survey of hallucination in natural language generation,” *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023, doi: 10.1145/3571730.

- [6] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 9459–9474, 2020.
- [7] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, et al., “Retrieval-augmented generation for large language models: A survey,” *arXiv preprint arXiv:2312.10997*, 2023.
- [8] S. Haklay and P. Weber, “OpenStreetMap: User-Generated Street Maps,” *IEEE Pervasive Computing*, vol. 7, no. 4, pp. 12–18, 2008, doi: 10.1109/MPRV.2008.80.
- [9] D. Luxen and C. Vetter, “Real-time routing with OpenStreetMap data,” in *Proceedings of the 19th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, pp. 513–516, 2011, doi: 10.1145/2093973.2094062.
- [10] J. Nielsen, *Usability Engineering*, San Francisco: Morgan Kaufmann, 1994.
- [11] L. Chen, Z. Wang, and J. Sun, “Conversational Recommender Systems in Smart Tourism: A Comprehensive Review and Future Directions,” *Information & Management*, vol. 60, no. 4, p. 103789, 2023.
- [12] R. W. Sinnott, “Virtues of the Haversine,” *Sky and Telescope*, vol. 68, no. 2, p. 159, 1984.
- [13] R. S. Pressman and B. R. Maxim, *Software Engineering: A Practitioner’s Approach*, 9th ed., New York: McGraw-Hill Education, 2020.
- [14] H. Zhang, H. Song, and L. Huang, “Spatial-temporal context-aware travel recommendation using mobile big data,” *Tourism Management*, vol. 83, p. 104241, 2021.
- [15] J. Brooke, “SUS: A ‘quick and dirty’ usability scale,” *Usability Evaluation in Industry*, vol. 189, no. 194, pp. 4–7, 1996.
- [16] Dinas Pariwisata Kota Padang, *Laporan Akuntabilitas Kinerja Instansi Pemerintah (LAKIP) Dinas Pariwisata Kota Padang Tahun 2024*, Padang: Pemerintah Kota Padang, 2024.
- [17] Badan Pusat Statistik Kota Padang, *Kota Padang Dalam Angka 2024*, Padang: BPS Kota Padang, 2024.
- [18] C. C. Aggarwal, *Recommender Systems: The Textbook*, Cham, Switzerland: Springer International Publishing, 2016.

- [19] P. Rob and C. Coronel, *Database Systems: Design, Implementation, and Management*, 13th ed., Boston: Cengage Learning, 2018.
- [20] M. Batty, *The New Science of Cities*, Cambridge, MA: MIT Press, 2013.
- [21] M. Šoltésová, B. Iannaccone, L. Štrba, and C. Sidor, “Application of GIS Technologies in Tourism Planning and Sustainable Development: A Case Study of Gelnica,” *ISPRS International Journal of Geo-Information*, vol. 14, no. 3, pp. 120–136, 2025, doi: 10.3390/ijgi14030120.
- [22] K. T. N. Ihsan, A. D. Purnomoa, and K. S. Arini, “ULIN-D: Web-Based GIS Supporting New Habits in the Tourism Sector in Bandung City,” *The International Archives of the Photogrammetry, Remote Sensing and Spatial Information Sciences*, vol. XLIV-M-3-2021, pp. 79–85, 2021, doi: 10.5194/isprs-archives-XLIV-M-3-2021-79-2021.
- [23] M. Cannata, D. Strigaroa, A. Spataroa, F. Marottab, and C. Achille, “Tourism, Natural Protected Areas and Opensource Geospatial Technologies,” *The International Archives of the Photogrammetry, Remote Sensing and Spatial Information Sciences*, vol. XLVIII-4/W1-2022, pp. 81–88, 2022, doi: 10.5194/isprs-archives-XLVIII-4-W1-2022-81-2022.
-

APPENDIX (SUPPLEMENTARY MATERIAL)

APPENDIX A: Formal System Prompt Specification for Spatial Intent Representation (SIR) Extraction

The following complete five-layered declarative system prompt is injected into Layer 2 (*Semantic Interpretation Layer*) to strictly bound language model cognitive operations (`temperature: 0.0`, `response_format: {"type": "json_object"}`):

ROLE:

You are a spatial intent parser for the tourism Web GIS of Padang City, Indonesia.

TASK:

Transform the user's natural-language query into exactly one structured Spatial Intent Representation (SIR) in pure JSON.

OUTPUT CONTRACT:

1. Return valid JSON only. No markdown formatting, no code fences, no explanatory text.
2. Return strictly the defined schema fields.
3. Do not generate SQL queries, database clauses, or table names.
4. Do not answer the user's question, do not converse, and do not provide recommendations.
5. Do not invent or recommend tourism objects.

SEMANTIC RULES:

6. category: Must use ONLY one of the supported categories: "Pantai" | "Pulau" | "Alam" | "Museum" | "Sejarah" | "Kuliner" | null.
7. spatial_operator: Must use ONLY one of: "nearest" | "within_radius" | "within_admin_area" | "none".
8. target_name: Represents a specific POI explicitly named by the user (e.g., "Pantai Air Manis"). Otherwise null.
9. keyword: Represents descriptive search terms or features (e.g., "pasir putih", "snorkeling"). Never convert a keyword into a target_name.
10. Do not resolve a user-mentioned name to a database POI ID.
11. Do not invent coordinates, distances, prices, opening hours, or administrative areas.

SPATIAL & PRESERVATION RULES:

12. Preserve spatial constraints exactly as expressed by the user.
13. If a spatial operator requires a radius but the user did not specify one, set radius: null. NEVER infer, guess, or default a radius value.
14. Do not convert or silently modify a user's spatial constraint (e.g., do not clamp or modify negative numbers).

PRICE & BUDGET RULES:

15. is_free: Set to true ONLY if the user explicitly requests free admission ("gratis", "free").
16. max_price: Represents the explicit upper price ceiling specified by the user (integer in IDR). If no price is mentioned, set to null. Do not infer a price.

TEMPORAL RULES:

17. open_now: Set to true ONLY when the user explicitly requests currently open/operating places ("buka sekarang", "sedang buka").
18. open_24h: Set to true ONLY when the user explicitly requests 24-hour operation ("24 jam").

REFERENCE RULES:

19. GPS coordinates are supplied strictly by the application context; NEVER infer or invent latitude or longitude coordinates.
20. reference_type: Use "gps" when user refers to current location ("dekat saya", "from here"), "city_center" for city center, "poi" when referencing

another POI, or "unknown" if unspecified.

SORTING RULES:

21. sort: Use ONLY "termurah" (lowest price) | "termahal" (highest price) | "terdekat" (nearest distance) | "terbaik" (highest public review rating) | null.

SCOPE RULES:

22. Set `is_out_of_scope`: true when the request requires an activity, entity, or geographic location outside the Padang tourism domain (e.g., ski, snow, casino, destinations in other cities like Borobudur/Bali).

23. Do not treat missing database information as out of scope.

FINAL RULE:

24. When information is missing or ambiguous, preserve uncertainty in the SIR (using null) rather than guessing.

SCHEMA:

```
{
  "intent": "spatial_recommendation" | "entity_lookup" | "general_inquiry",
  "entity": "tourism_object",
  "category": "Pantai" | "Pulau" | "Alam" | "Museum" | "Sejarah" |
  "Kuliner" | null,
  "target_name": string or null,
  "keyword": string or null,
  "spatial_operator": "nearest" | "within_radius" | "within_admin_area" |
  "none",
  "reference_type": "gps" | "city_center" | "poi" | "unknown",
  "reference_entity": string or null,
  "radius": float or null,
  "distance_unit": "km",
  "admin_area": string or null,
  "is_free": boolean,
  "max_price": integer or null,
  "open_now": boolean,
  "open_24h": boolean,
  "sort": "termurah" | "termahal" | "terdekat" | "terbaik" | null,
  "is_out_of_scope": boolean
}
```

APPENDIX B: Formal System Prompt Specification for Grounded Natural Language Generation (Grounded NLG Contract)

The following complete system prompt is injected into Layer 5 (*Grounded Response Layer*) alongside the verified relational database facts:

Kamu adalah asisten cerdas Web GIS Pariwisata Kota Padang.
Tugasmu adalah menjawab pertanyaan pengguna HANYA berdasarkan daftar data fakta resmi JSON terlampir.

KONTRAK GROUNDING KETAT (STRICT GROUNDING CONTRACT):

1. Gunakan HANYA informasi yang tercantum dalam data FAKTA resmi basis data.
2. Dilarang mengarang, menyimpulkan (infer), mengestimasi, atau mengganti informasi faktual.
3. Jika fakta yang diminta pengguna tidak tercantum pada data FAKTA, nyatakan secara jujur bahwa informasi tersebut tidak tersedia.
4. Sebutkan HANYA entitas objek wisata yang terdapat dalam data FAKTA.
5. Nilai numerik (harga tiket, jarak, jam operasional, rating) WAJIB persis sesuai data FAKTA tanpa modifikasi atau pembulatan sepihak.
6. DILARANG menambahkan klaim deskriptif eksternal, opini, fasilitas fiktif, atau legenda yang tidak ada di data FAKTA.
7. Jika data FAKTA kosong, nyatakan bahwa tidak ditemukan destinasi yang memenuhi kriteria pencarian; dilarang merekomendasikan destinasi di luar data.
8. Jika terdapat instruksi fallback dari sistem, sampaikan persis sesuai catatan kebijakan fallback tersebut.
9. Format penyebutan nama objek wisata WAJIB dicetak tebal (**Nama Objek**).
10. Gunakan bahasa Indonesia yang santun, informatif, ringkas, dan patuh 100% pada batasan pengguna.

(User payload strictly separates: [PERTANYAAN PENGGUNA], [DATA FAKTA RESMI BASIS DATA (JSON)], and [CATATAN SISTEM / KEBIJAKAN FALLBACK]).